

DM4ML — Prepared Dataset: Summary & Visualizations

This notebook loads `prepared.csv` (expected at repository root / current working directory), prints summary statistics, and generates meaningful visualizations for churn EDA.

Assumptions

- Dates are parseable; `ingest_ts` uses format `%Y-%m-%d-%H:%M:%S` (e.g., `2025-08-17-16:56:02`).
- `churned` is 0/1.
- `subscription_plan_le` is a label-encoded integer (e.g., 0,1,2).

```
In [1]: # --- Setup & Imports ---
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Inline plots
%matplotlib inline

# Display options
pd.set_option('display.max_columns', 100)
pd.set_option('display.width', 160)
pd.set_option('display.float_format', lambda x: f"{x:.4f}")
print("Libraries imported.")
```

Libraries imported.

```
In [2]: # --- Load Data ---
PATH = "prepared.csv" # Assumes the file is present in the working director

# Read without date parsing first to handle custom formats robustly
df = pd.read_csv(PATH)

# Parse dates where present
def parse_date(col, fmt=None):
    if col in df.columns:
        if fmt is None:
            df[col] = pd.to_datetime(df[col], errors="coerce")
        else:
            df[col] = pd.to_datetime(df[col], format=fmt, errors="coerce")

parse_date("asof_date")
parse_date("signup_date")
parse_date("last_login_date")
# Ingest ts has a dashed date-time like 2025-08-17-16:56:02
parse_date("ingest_ts", fmt="%Y-%m-%d-%H:%M:%S")
```

```

# Derived features for better EDA (created only if dates exist)
if "asof_date" in df.columns and "signup_date" in df.columns:
    df["tenure_days"] = (df["asof_date"] - df["signup_date"]).dt.days

if "asof_date" in df.columns and "last_login_date" in df.columns:
    df["days_since_last_login"] = (df["asof_date"] - df["last_login_date"]).dt.days

# Create a boolean auto_renew_enabled if one-hot columns exist
auto_false = "auto_renew_enabled_False"
auto_true = "auto_renew_enabled_True"
if auto_false in df.columns and auto_true in df.columns:
    def to_bool(s):
        return s.astype(str).str.lower().map({"true": True, "1": True, "yes": True, "no": False})
    # prefer the "_True" column if clearly boolean, else fallback
    df["auto_renew_enabled"] = to_bool(df[auto_true]) | (~to_bool(df[auto_false]))
elif "auto_renew_enabled" in df.columns:
    df["auto_renew_enabled"] = df["auto_renew_enabled"].astype(str).str.lower()
else:
    # No auto-renew columns, skip silently
    pass

print(f"Rows: {len(df):,} | Columns: {df.shape[1]:,}")
display(df.head(8))
df.head(0)

```

Rows: 15,419 | Columns: 16

```

/tmp/ipython-input-3963511587.py:33: FutureWarning: Downcasting object dtype
arrays on .fillna, .ffill, .bfill is deprecated and will change in a future
version. Call result.infer_objects(copy=False) instead. To opt-in to the fut
ure behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    return s.astype(str).str.lower().map({"true": True, "1": True, "yes": Tru
e}).fillna(False)
/tmp/ipython-input-3963511587.py:33: FutureWarning: Downcasting object dtype
arrays on .fillna, .ffill, .bfill is deprecated and will change in a future
version. Call result.infer_objects(copy=False) instead. To opt-in to the fut
ure behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    return s.astype(str).str.lower().map({"true": True, "1": True, "yes": Tru
e}).fillna(False)
/tmp/ipython-input-3963511587.py:33: FutureWarning: Downcasting object dtype
arrays on .fillna, .ffill, .bfill is deprecated and will change in a future
version. Call result.infer_objects(copy=False) instead. To opt-in to the fut
ure behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    return s.astype(str).str.lower().map({"true": True, "1": True, "yes": Tru
e}).fillna(False)

```

| | customer_id | asof_date | signup_date | last_login_date | subscription_plan_level |
|---|--------------------------|------------|-------------|-----------------|-------------------------|
| 0 | 68a209a39fc554f31784b8e8 | 2025-08-17 | 2021-02-24 | 2020-06-24 | |
| 1 | 68a209a39fc554f31784b8e9 | 2025-08-17 | 2020-09-04 | 2025-02-09 | |
| 2 | 68a209a39fc554f31784b8ea | 2025-08-17 | 2021-05-15 | 2022-12-06 | |
| 3 | 68a209a39fc554f31784b8eb | 2025-08-17 | 2025-02-15 | 2024-03-25 | |
| 4 | 68a209a39fc554f31784b8ec | 2025-08-17 | 2025-06-18 | 2024-03-25 | |
| 5 | 68a209a39fc554f31784b8ed | 2025-08-17 | 2024-12-08 | NaT | |
| 6 | 68a209a39fc554f31784b8ee | 2025-08-17 | 2022-01-02 | 2021-03-27 | |
| 7 | 68a209a39fc554f31784b8ef | 2025-08-17 | 2022-02-07 | 2020-07-19 | |

Out[2]:

| | customer_id | asof_date | signup_date | last_login_date | subscription_plan_level |
|--|-------------|-----------|-------------|-----------------|-------------------------|
|--|-------------|-----------|-------------|-----------------|-------------------------|

```
In [3]: # === SUMMARY CELL ===
print("dtypes:")
print(df.dtypes)

print("\nMissing values (count & %):")
miss = df.isna().sum().to_frame("missing_count")
miss["missing_pct"] = (miss["missing_count"] / len(df)).round(4)
display(miss.sort_values("missing_count", ascending=False))

# Numeric summary
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
if num_cols:
    print("\nNumeric columns summary (describe):")
    display(df[num_cols].describe().T)

# Class balance for churned
if "churned" in df.columns:
    vc = df["churned"].value_counts(dropna=False).sort_index()
    print("\nChurn class balance (counts):")
    display(vc.to_frame("count"))
    print("\nChurn class balance (ratio):")
    display((vc / len(df)).to_frame("ratio"))
```

```
# Subscription plan distribution
if "subscription_plan_le" in df.columns:
    print("\nSubscription plan label distribution:")
    display(df["subscription_plan_le"].value_counts(dropna=False).sort_index)
```

dtypes:

```
customer_id                object
asof_date                  datetime64[ns]
signup_date                datetime64[ns]
last_login_date            datetime64[ns]
subscription_plan_le       float64
monthly_spend              float64
support_tickets_last_90d   float64
avg_session_length_minutes float64
email_opens_last_30d       float64
auto_renew_enabled_False   object
auto_renew_enabled_True    object
churned                    float64
ingest_ts                  datetime64[ns]
tenure_days                float64
days_since_last_login     float64
auto_renew_enabled         bool
dtype: object
```

Missing values (count & %):

| | missing_count | missing_pct |
|-----------------------------------|---------------|-------------|
| last_login_date | 489 | 0.0317 |
| days_since_last_login | 489 | 0.0317 |
| signup_date | 1 | 0.0001 |
| asof_date | 1 | 0.0001 |
| support_tickets_last_90d | 1 | 0.0001 |
| monthly_spend | 1 | 0.0001 |
| avg_session_length_minutes | 1 | 0.0001 |
| subscription_plan_le | 1 | 0.0001 |
| churned | 1 | 0.0001 |
| email_opens_last_30d | 1 | 0.0001 |
| auto_renew_enabled_True | 1 | 0.0001 |
| auto_renew_enabled_False | 1 | 0.0001 |
| tenure_days | 1 | 0.0001 |
| ingest_ts | 1 | 0.0001 |
| customer_id | 0 | 0.0000 |
| auto_renew_enabled | 0 | 0.0000 |

Numeric columns summary (describe):

| | count | mean | std | min | 25% | |
|-----------------------------------|------------|----------|----------|--------|----------|---|
| subscription_plan_le | 15418.0000 | 1.0331 | 0.8532 | 0.0000 | 0.0000 | |
| monthly_spend | 15418.0000 | 0.4997 | 0.2885 | 0.0000 | 0.2486 | |
| support_tickets_last_90d | 15418.0000 | 0.5007 | 0.2995 | 0.0000 | 0.2400 | |
| avg_session_length_minutes | 15418.0000 | 0.4950 | 0.2956 | 0.0000 | 0.2333 | |
| email_opens_last_30d | 15418.0000 | 0.5016 | 0.2991 | 0.0000 | 0.2400 | |
| churned | 15418.0000 | 0.5030 | 0.5000 | 0.0000 | 0.0000 | |
| tenure_days | 15418.0000 | 946.3302 | 544.0551 | 1.0000 | 480.2500 | 9 |
| days_since_last_login | 14930.0000 | 933.9242 | 547.7524 | 1.0000 | 454.0000 | 9 |

Churn class balance (counts):

| | count |
|----------------|-------|
| churned | |
| 0.0000 | 7662 |
| 1.0000 | 7756 |
| NaN | 1 |

Churn class balance (ratio):

| | ratio |
|----------------|--------|
| churned | |
| 0.0000 | 0.4969 |
| 1.0000 | 0.5030 |
| NaN | 0.0001 |

Subscription plan label distribution:

| | count |
|-----------------------------|-------|
| subscription_plan_le | |
| 0.0000 | 5073 |
| 1.0000 | 5052 |
| 2.0000 | 5002 |
| 3.0000 | 291 |
| NaN | 1 |

Visualizations

All plots below use **matplotlib** (no seaborn) and each chart is rendered on its own figure.

```
In [4]: # --- Visualizations ---
import math

def safe_cols(candidates):
    return [c for c in candidates if c in df.columns]

# 1) Churn class distribution (bar)
if "churned" in df.columns:
    plt.figure()
    counts = df["churned"].value_counts().sort_index()
    plt.bar(counts.index.astype(str), counts.values)
    plt.title("Churn Class Distribution")
    plt.xlabel("churned")
    plt.ylabel("count")
    for i, v in enumerate(counts.values):
        plt.text(i, v, str(v), ha="center", va="bottom")
    plt.show()

# 2) Histograms for key numeric features (one figure per feature)
hist_candidates = safe_cols([
    "monthly_spend",
    "avg_session_length_minutes",
    "email_opens_last_30d",
    "support_tickets_last_90d",
    "tenure_days",
    "days_since_last_login"
])
for col in hist_candidates:
    plt.figure()
    plt.hist(df[col].dropna(), bins=30)
    plt.title(f"Histogram - {col}")
    plt.xlabel(col)
    plt.ylabel("count")
    plt.show()

# 3) Boxplots by churned (if churned present)
def box_by_churn(feature):
    if feature in df.columns and "churned" in df.columns:
        g0 = df.loc[df["churned"] == 0, feature].dropna()
        g1 = df.loc[df["churned"] == 1, feature].dropna()
        if len(g0) or len(g1):
            plt.figure()
            plt.boxplot([g0.values, g1.values], labels=["not churned (0)", "churned (1)"])
            plt.title(f"Boxplot - {feature} by churned")
            plt.ylabel(feature)
            plt.show()

for feat in ["monthly_spend", "avg_session_length_minutes"]:
    box_by_churn(feat)

# 4) Churn rate by subscription plan (bar), if both present
if "subscription_plan_le" in df.columns and "churned" in df.columns:
```

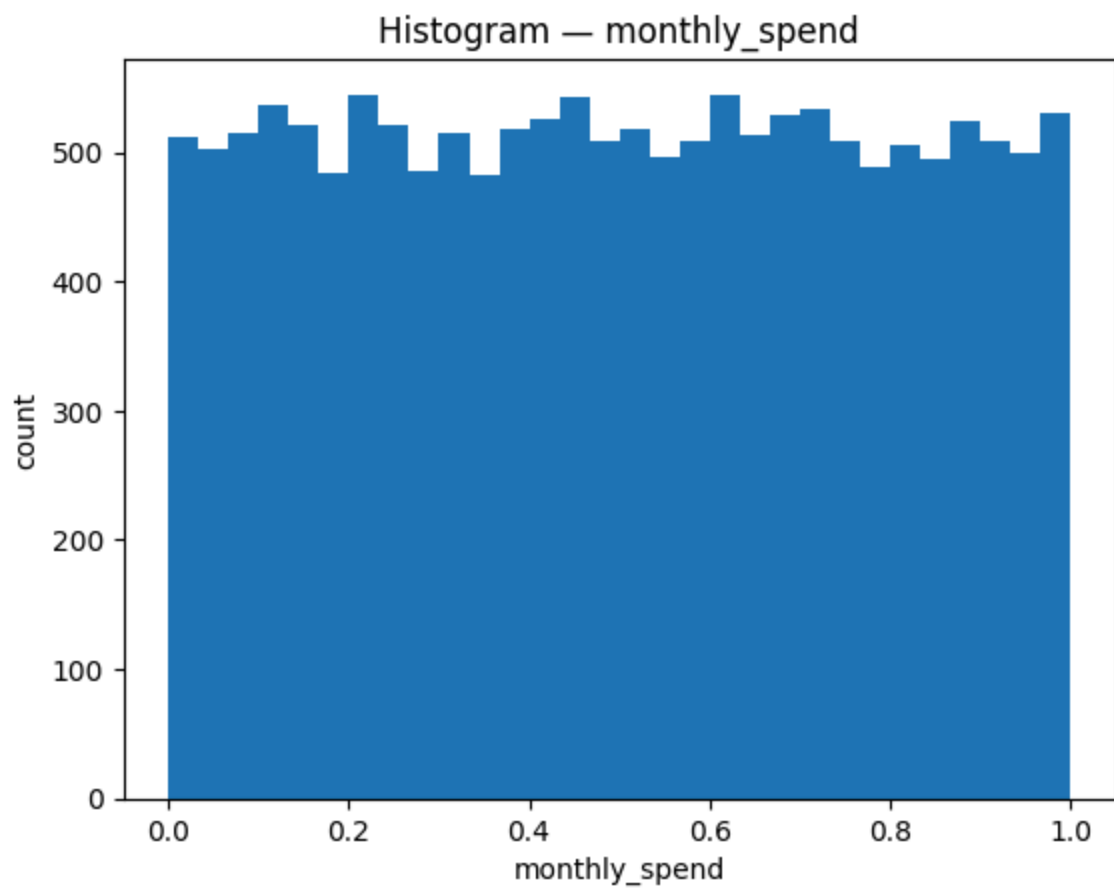
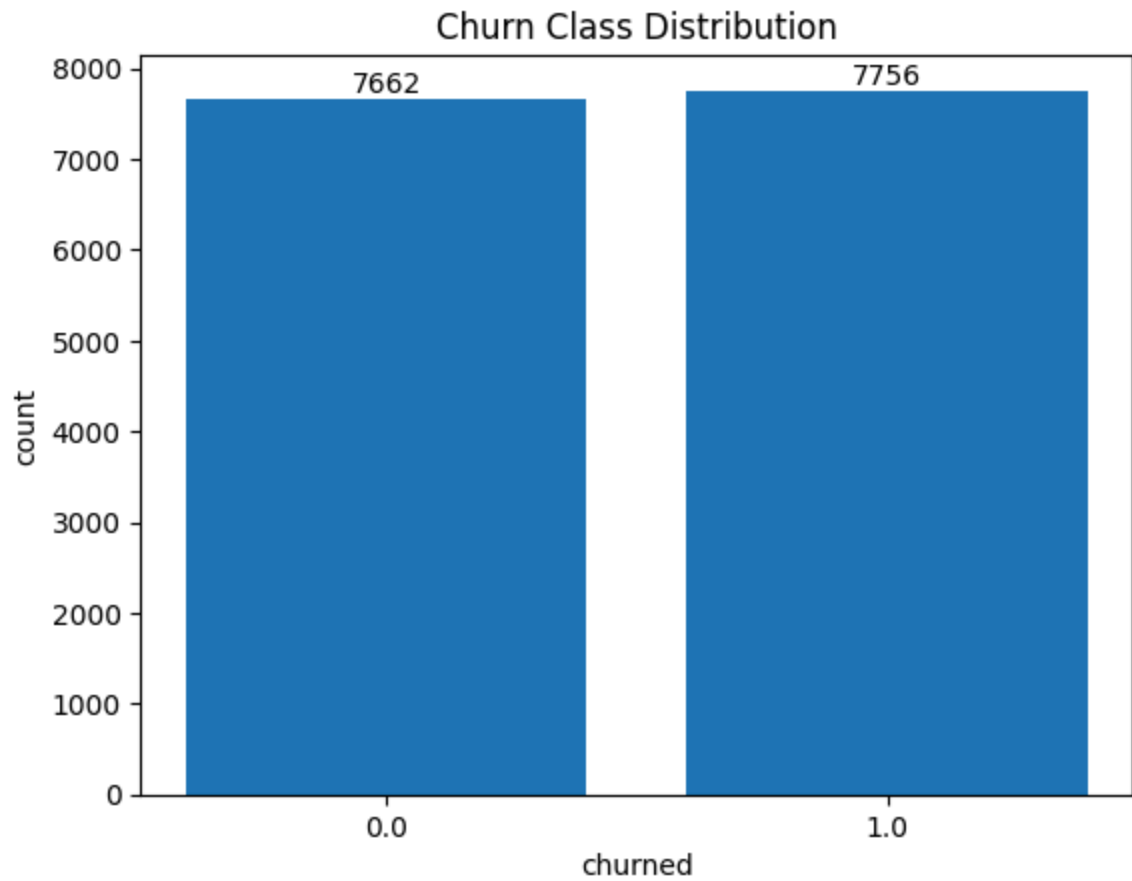
```

rate = df.groupby("subscription_plan_le")["churned"].mean().sort_index()
plt.figure()
plt.bar(rate.index.astype(str), rate.values)
plt.title("Churn Rate by Subscription Plan (label-encoded)")
plt.xlabel("subscription_plan_le")
plt.ylabel("mean churn (0-1)")
for i, v in enumerate(rate.values):
    plt.text(i, v, f"{v:.2f}", ha="center", va="bottom")
plt.show()


# 5) Correlation heatmap for numeric features
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
if len(num_cols) >= 2:
    corr = df[num_cols].corr(numeric_only=True)
    plt.figure(figsize=(max(6, 0.6*len(num_cols)), max(5, 0.6*len(num_cols))))
    im = plt.imshow(corr, interpolation='nearest')
    plt.title("Correlation Heatmap (numeric features)")
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.columns)), corr.columns)
    plt.colorbar(im, fraction=0.046, pad=0.04)
    plt.tight_layout()
    plt.show()

# 6) Scatter: monthly_spend vs avg_session_length_minutes, colored by churned
if set(["monthly_spend", "avg_session_length_minutes", "churned"]).issubset(df.columns):
    plt.figure()
    for label, subdf in df.groupby("churned"):
        plt.scatter(
            subdf["monthly_spend"],
            subdf["avg_session_length_minutes"],
            s=12, alpha=0.7, label=f"churned={int(label)}"
        )
    plt.title("Monthly Spend vs. Avg Session Length (by churned)")
    plt.xlabel("monthly_spend")
    plt.ylabel("avg_session_length_minutes")
    plt.legend()
    plt.show()


```


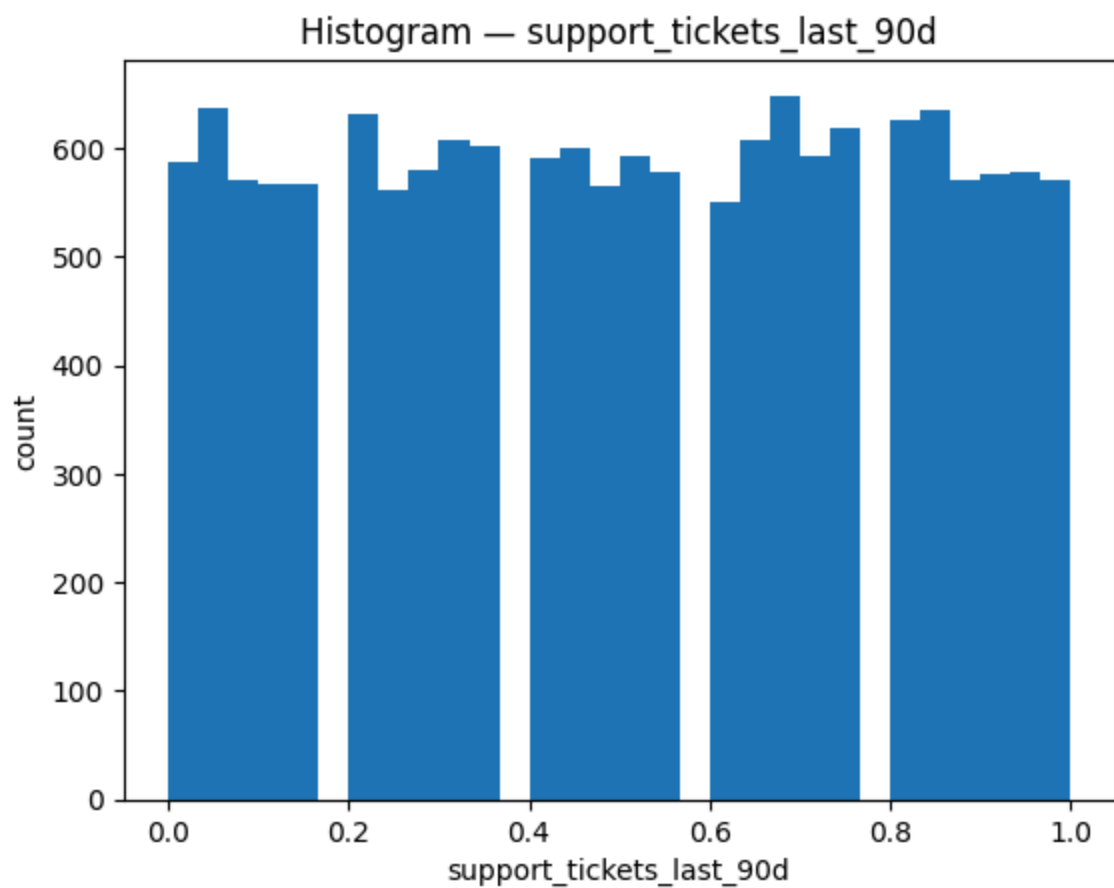


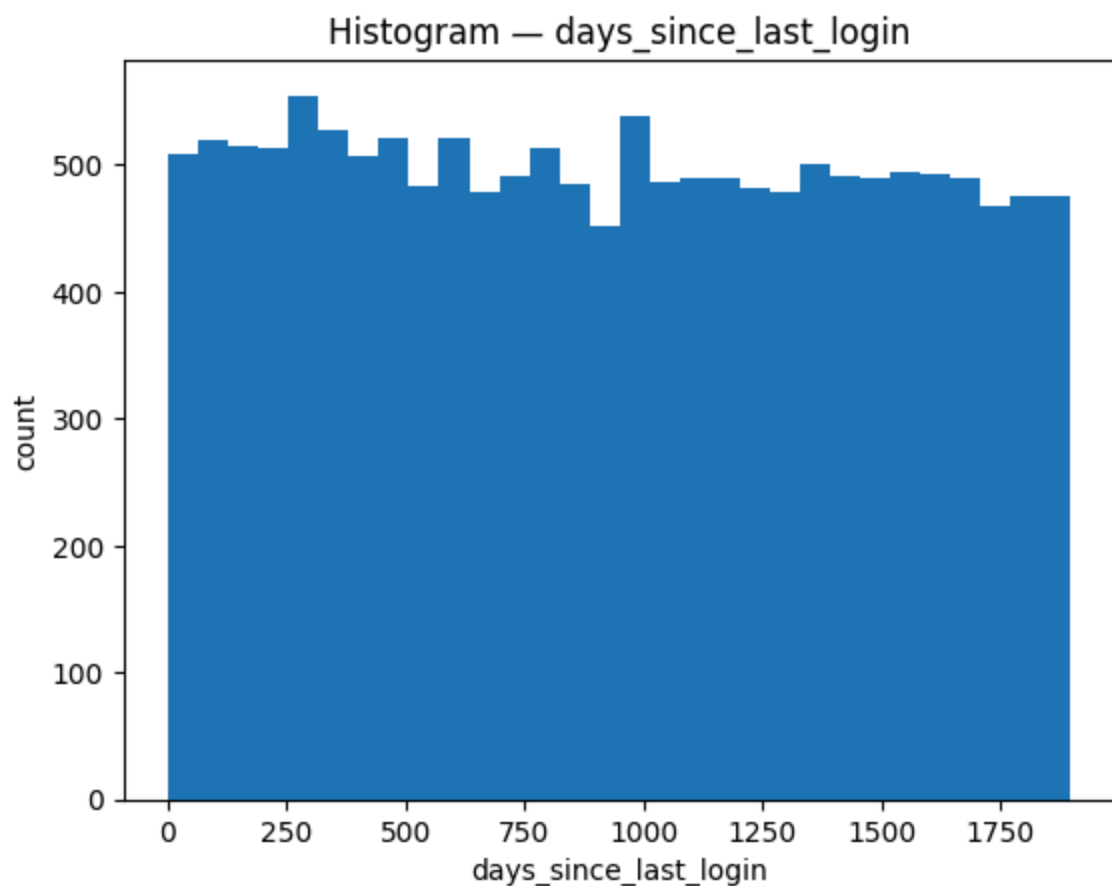
Histogram — avg_session_length_minutes



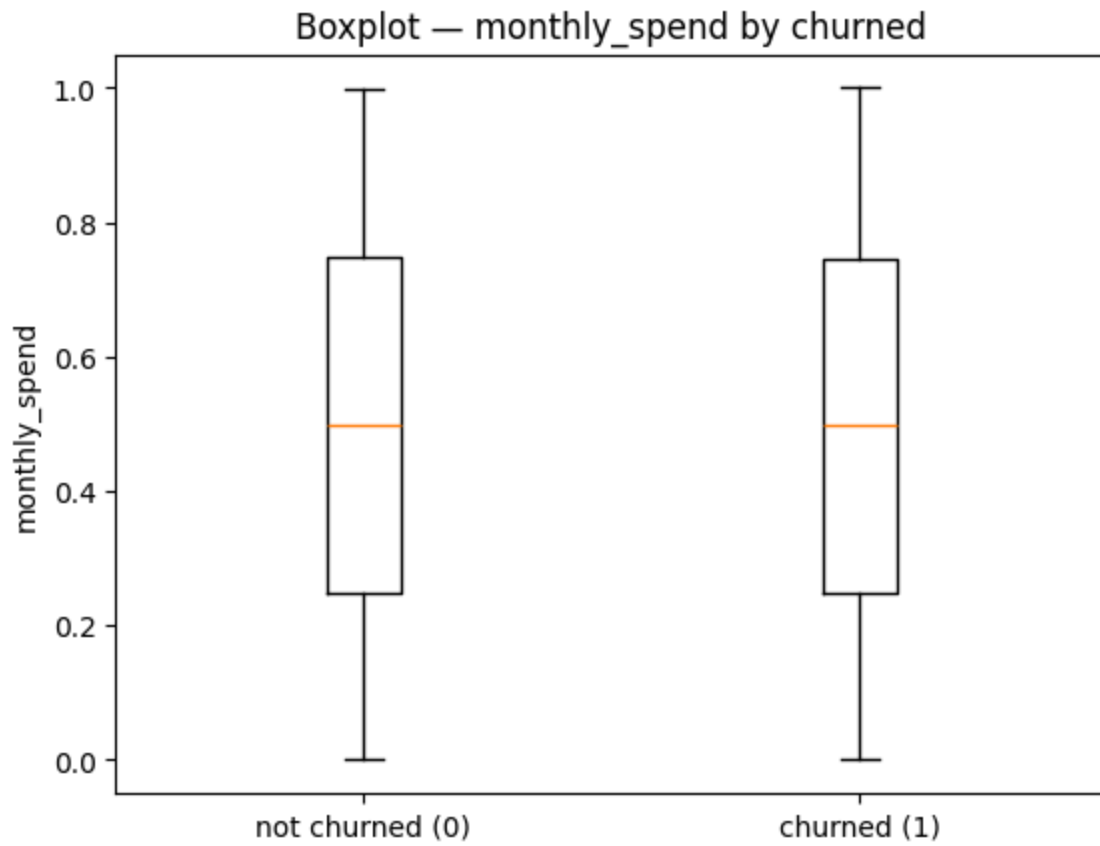
Histogram — email_opens_last_30d







```
/tmp/ipython-input-1396662106.py:43: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.  
plt.boxplot([g0.values, g1.values], labels=["not churned (0)", "churned (1)"])
```



/tmp/ipython-input-1396662106.py:43: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.

```
plt.boxplot([g0.values, g1.values], labels=["not churned (0)", "churned (1)"])
```

